

Exam 2 review

Questions

Table of contents

| | |
|--|-----------|
| Blizzard salaries | 2 |
| Question 1 | 3 |
| Question 2 | 3 |
| Question 3 | 4 |
| Question 4 | 4 |
| Question 5 | 5 |
| Question 6 | 7 |
| | |
| Movies | 8 |
| Score vs. runtime | 9 |
| Question 7 | 11 |
| Question 8 | 11 |
| Question 9 | 12 |
| Score vs. runtime or year | 12 |
| Question 10 | 13 |
| Question 11 | 13 |
| Score vs. runtime and rating | 14 |
| Question 12 | 14 |
| Question 13 | 15 |
| Question 14 | 15 |
| Question 15 | 16 |
| | |
| Miscellaneous | 16 |
| Question 16 | 16 |
| Question 17 | 16 |
| | |
| Building a spam filter | 17 |
| Question 18 | 17 |
| Question 19 | 18 |

Note

Suggested answers can be found [here](#), but resist the urge to peek before you go through it yourself.

Blizzard salaries

In 2020, employees of Blizzard Entertainment circulated a spreadsheet to anonymously share salaries and recent pay increases amidst rising tension in the video game industry over wage disparities and executive compensation. (Source: [Blizzard Workers Share Salaries in Revolt Over Pay](#))

The name of the data frame used for this analysis is `blizzard_salary` and the variables are:

- **percent_incr**: Raise given in July 2020, as percent increase with values ranging from 1 (1% increase) to 21.5 (21.5% increase)
 - **salary_type**: Type of salary, with levels **Hourly** and **Salaried**
 - **annual_salary**: Annual salary, in USD, with values ranging from \$50,939 to \$216,856.
 - **performance_rating**: Most recent review performance rating, with levels **Poor**, **Successful**, **High**, and **Top**. The **Poor** level is the lowest rating and the **Top** level is the highest rating.

The first ten rows of `blizzard_salary` are shown below:

```
# A tibble: 409 x 4
  percent_incr salary_type annual_salary performance_rating
  <dbl> <chr>           <dbl> <chr>
1 1     Salaried          1   High
2 1     Salaried          1   Successful
3 1     Salaried          1   High
4 1     Hourly            33987. Successful
5 NA    Hourly            34798. High
6 NA    Hourly            35360  <NA>
7 NA    Hourly            37440  <NA>
8 0     Hourly            37814. <NA>
9 4     Hourly            41101. Top
10 1.2   Hourly          42328  <NA>
# i 399 more rows
```

Question 1

You fit a model for predicting raises (`percent_incr`) from salaries (`annual_salary`). We'll call this model `raise_1_fit`. A tidy output of the model is shown below.

```
# A tibble: 2 x 5
  term      estimate std.error statistic p.value
  <chr>     <dbl>     <dbl>     <dbl>     <dbl>
1 (Intercept) 1.87      0.432      4.33  0.0000194
2 annual_salary 0.0000155 0.00000452    3.43  0.000669
```

Which of the following is the best interpretation of the slope coefficient?

- a. For every additional \$1,000 of annual salary, the model predicts the raise to be higher, on average, by 1.55%.
- b. For every additional \$1,000 of annual salary, the raise goes up by 0.0155%.
- c. For every additional \$1,000 of annual salary, the model predicts the raise to be higher, on average, by 0.0155%.
- d. For every additional \$1,000 of annual salary, the model predicts the raise to be higher, on average, by 1.87%.

Question 2

You then fit a model for predicting raises (`percent_incr`) from salaries (`annual_salary`) and performance ratings (`performance_rating`). We'll call this model `raise_2_fit`. Which of the following is definitely true based on the information you have so far?

- a. Intercept of `raise_2_fit` is higher than intercept of `raise_1_fit`.
- b. Slope of `raise_2_fit` is higher than RMSE of `raise_1_fit`.
- c. Adjusted R^2 of `raise_2_fit` is higher than adjusted R^2 of `raise_1_fit`.
- d. R^2 of `raise_2_fit` is higher R^2 of `raise_1_fit`.

Question 3

The tidy model output for the `raise_2_fit` model you fit is shown below.

| # A tibble: 5 x 5 | term | estimate | std.error | statistic | p.value |
|-------------------|------------------------------|------------|------------|-----------|----------|
| | <chr> | <dbl> | <dbl> | <dbl> | <dbl> |
| 1 | (Intercept) | 3.55 | 0.508 | 6.99 | 1.99e-11 |
| 2 | annual_salary | 0.00000989 | 0.00000436 | 2.27 | 2.42e- 2 |
| 3 | performance_ratingPoor | -4.06 | 1.42 | -2.86 | 4.58e- 3 |
| 4 | performance_ratingSuccessful | -2.40 | 0.397 | -6.05 | 4.68e- 9 |
| 5 | performance_ratingTop | 2.99 | 0.715 | 4.18 | 3.92e- 5 |

When your teammate sees this model output, they remark “The coefficient for `performance_ratingSuccessful` is negative. That’s weird. I guess it means that people who get successful performance ratings get lower raises.” How would you respond to your teammate?

Question 4

Ultimately, your teammate decides they don’t like the negative slope coefficients in the model output you created (not that there’s anything wrong with negative slope coefficients!), does something else, and comes up with the following model output.

| # A tibble: 5 x 5 | term | estimate | std.error | statistic | p.value |
|-------------------|------------------------------|------------|------------|-----------|------------|
| | <chr> | <dbl> | <dbl> | <dbl> | <dbl> |
| 1 | (Intercept) | -0.511 | 1.47 | -0.347 | 0.729 |
| 2 | annual_salary | 0.00000989 | 0.00000436 | 2.27 | 0.0242 |
| 3 | performance_ratingSuccessful | 1.66 | 1.42 | 1.17 | 0.242 |
| 4 | performance_ratingHigh | 4.06 | 1.42 | 2.86 | 0.00458 |
| 5 | performance_ratingTop | 7.05 | 1.53 | 4.60 | 0.00000644 |

Unfortunately they didn’t write their code in a Quarto document, instead just wrote some code in the Console and then lost track of their work. They remember using the `fct_relevel()` function and doing something like the following:

```
blizzard_salary <- blizzard_salary |>  
  mutate(performance_rating = fct_relevel(performance_rating, ___))
```

What should they put in the blanks to get the same model output as above?

- a. “Poor”, “Successful”, “High”, “Top”
- b. “Successful”, “High”, “Top”
- c. “Top”, “High”, “Successful”, “Poor”
- d. Poor, Successful, High, Top

Question 5

Suppose we fit a model to predict `percent_incr` from `annual_salary` and `salary_type`. A tidy output of the model is shown below.

```
# A tibble: 3 x 5  
  term            estimate std.error statistic p.value  
  <chr>          <dbl>     <dbl>      <dbl>    <dbl>  
1 (Intercept)    1.24      0.570      2.18    0.0300  
2 annual_salary  0.0000137 0.00000464   2.96   0.00329  
3 salary_typeSalaried 0.913     0.544      1.68   0.0938
```

Which of the following visualizations represent this model? Explain your reasoning.




Figure 1: Visualizations of the relationship between percent increase, annual salary, and salary type

Question 6

Suppose you now fit a model to predict the natural log of percent increase, `log(percent_incr)`, from performance rating. The model is called `raise_4_fit`.

You're provided the following:

```
tidy(raise_4_fit) |>
  select(term, estimate) |>
  mutate(exp_estimate = exp(estimate))
```

```
# A tibble: 4 x 3
  term                estimate exp_estimate
  <chr>              <dbl>      <dbl>
1 (Intercept)        -7.15      0.000786
2 performance_ratingSuccessful    6.93     1025.
3 performance_ratingHigh         8.17     3534.
4 performance_ratingTop          8.91     7438.
```

Based on this, which of the following is true?

- a. The model predicts that the percentage increase employees with Successful performance get, on average, is higher by 10.25% compared to the employees with Poor performance rating.
- b. The model predicts that the percentage increase employees with Successful performance get, on average, is higher by 6.93% compared to the employees with Poor performance rating.
- c. The model predicts that the percentage increase employees with Successful performance get, on average, is higher by a factor of 1025 compared to the employees with Poor performance rating.
- d. The model predicts that the percentage increase employees with Successful performance get, on average, is higher by a factor of 6.93 compared to the employees with Poor performance rating.

Movies

The data for this part comes from the Internet Movie Database (IMDB). Specifically, the data are a random sample of movies released between 1980 and 2020.

The name of the data frame used for this analysis is `movies`, and it contains the variables shown in Table 1.

Table 1: Data dictionary for `movies`

| Variable | Description |
|------------------------------|---|
| <code>name</code> | name of the movie |
| <code>rating</code> | rating of the movie (R, PG, etc.) |
| <code>genre</code> | main genre of the movie. |
| <code>runtime</code> | duration of the movie |
| <code>year</code> | year of release |
| <code>release_date</code> | release date (YYYY-MM-DD) |
| <code>release_country</code> | release country |
| <code>score</code> | IMDB user rating |
| <code>votes</code> | number of user votes |
| <code>director</code> | the director |
| <code>writer</code> | writer of the movie |
| <code>star</code> | main actor/actress |
| <code>country</code> | country of origin |
| <code>budget</code> | the budget of a movie (some movies don't have this, so it appears as 0) |
| <code>gross</code> | revenue of the movie |
| <code>company</code> | the production company |

The first thirty rows of the `movies` data frame are shown in Table 2, with variable types suppressed (since we'll ask about them later).

Score vs. runtime

In this part, we fit a model predicting `score` from `runtime` and name it `score_runtime_fit`.

```
score_runtime_fit <- linear_reg() |>  
  fit(score ~ runtime, data = movies)
```

Figure 2 visualizes the relationship between `score` and `runtime` as well as the linear model for predicting `score` from `runtime`. The first three movies in Table 2 are labeled in the visualization as well. Answer all questions in this part based on Figure 2.




Figure 2: Scatterplot of `score` vs. `runtime` for `movies`.

Table 2

First 30 rows of `movies`, with variable types suppressed.

A tibble: 500 x 16

| | name | score | runtime | genre | rating | release_country | release_date |
|----|----------------|-----------|---------|-----------|------------------|-----------------|---------------|
| 1 | Blue City | 4.4 | 83 mins | Action | R | United States | 1986-05-02 |
| 2 | Winter Sleep | 8.1 | 196 | Drama | Not Rated | Turkey | 2014-06-12 |
| 3 | Rang De Basan~ | 8.1 | 167 | Comedy | Not Rated | United States | 2006-01-26 |
| 4 | Pokémon Detec~ | 6.6 | 104 | Action | PG | United States | 2019-05-10 |
| 5 | A Bad Moms Ch~ | 5.6 | 104 | Comedy | R | United States | 2017-11-01 |
| 6 | Replicas | 5.5 | 107 | Drama | PG-13 | United States | 2019-01-11 |
| 7 | Windy City | 5.8 | 103 | Drama | R | Uruguay | 1986-01-01 |
| 8 | War for the P~ | 7.4 | 140 | Action | PG-13 | United States | 2017-07-14 |
| 9 | Tales from th~ | 6.4 | 98 | Crime | R | United States | 1995-05-24 |
| 10 | Fire with Fire | 6.5 | 103 | Drama | PG-13 | United States | 1986-05-09 |
| 11 | Raising Helen | 6 | 119 | Comedy | PG-13 | United States | 2004-05-28 |
| 12 | Feeling Minne~ | 5.4 | 99 | Comedy | R | United States | 1996-09-13 |
| 13 | The Babe | 5.9 | 115 | Biography | PG | United States | 1992-04-17 |
| 14 | The Real Blon~ | 6 | 105 | Comedy | R | United States | 1998-02-27 |
| 15 | To vlemma tou~ | 7.6 | 176 | Drama | Not Rated | United States | 1997-11-01 |
| 16 | Going the Dis~ | 6.3 | 102 | Comedy | R | United States | 2010-09-03 |
| 17 | Jung on zo | 6.8 | 103 | Action | R | Hong Kong | 1993-06-24 |
| 18 | Rita, Sue and~ | 6.5 | 93 | Comedy | R | United Kingdom | 1987-05-29 |
| 19 | Phone Booth | 7 | 81 | Crime | R | United States | 2003-04-04 |
| 20 | Happy Death D~ | 6.6 | 96 | Comedy | PG-13 | United States | 2017-10-13 |
| 21 | Barely Legal | 4.7 | 90 | Comedy | R | Thailand | 2006-05-25 |
| 22 | Three Kings | 7.1 | 114 | Action | R | United States | 1999-10-01 |
| 23 | Menace II Soc~ | 7.5 | 97 | Crime | R | United States | 1993-05-26 |
| 24 | Four Rooms | 6.8 | 98 | Comedy | R | United States | 1995-12-25 |
| 25 | Quartet | 6.8 | 98 | Comedy | PG-13 | United States | 2013-03-01 |
| 26 | Tape | 7.2 | 86 | Drama | R | Denmark | 2002-07-12 |
| 27 | Marked for De~ | 6 | 93 | Action | R | United States | 1990-10-05 |
| 28 | Congo | 5.2 | 109 | Action | PG-13 | United States | 1995-06-09 |
| 29 | Stop-Loss | 6.4 | 112 | Drama | R | United States | 2008-03-28 |
| 30 | Con Air | 6.9 | 115 | Action | R | United States | 1997-06-06 |
| | budget | gross | votes | year | director | writer | star |
| 1 | 10000000 | 6947787 | 1100 | 1986 | Michelle Manning | Ross Macdonal~ | Judd Nelson |
| 2 | NA | 4018705 | 48000 | 2014 | Nuri Bilge Ceyl~ | Ebru Ceylan | Haluk Bilgin~ |
| 3 | NA | 10800778 | 115000 | 2006 | Rakeysh Ompraka~ | Renzil D'Sil~ | Aamir Khan |
| 4 | 150000000 | 433921300 | 146000 | 2019 | Rob Letterman | Dan Hernandez | Ryan Reynolds |
| 5 | 28000000 | 130560428 | 46000 | 2017 | Jon Lucas | Jon Lucas | Mila Kunis |
| 6 | 30000000 | 9330075 | 34000 | 2018 | Jeffrey Nachman~ | Chad St. John | Keanu Reeves |
| 7 | NA | 343890 | 262 | 1984 | Armyan Bernstein | Armyan Berns~ | John Shea |
| 8 | 150000000 | 490719763 | 235000 | 2017 | Matt Reeves | Mark Bomback | Andy Serkis |
| 9 | 6000000 | 11837928 | 7400 | 1995 | Rusty Cundieff | Rusty Cundie~ | Clarence Wil~ |
| 10 | NA | 4636169 | 1500 | 1986 | Duncan Gibbins | Bill Phillips | Craig Sheffer |
| 11 | 50000000 | 49718611 | 36000 | 2004 | Garry Marshall | Patrick J. C~ | Kate Hudson |
| 12 | NA | 3124440 | 11000 | 1996 | Steven Baigelman | Steven Baige~ | Keanu Reeves |
| 13 | NA | 19930973 | 9300 | 1992 | Arthur Hiller | John Fusco | John Goodman |
| 14 | NA | 83488 | 3900 | 1997 | Tom DiCillo | Tom DiCillo | Matthew Modi~ |
| 15 | NA | NA | 6400 | 1995 | Theodoros Angel~ | Theodoros An~ | Harvey Keitel |

Question 7

Partial code for producing Figure 2 is given below. Which of the following goes in the blank on Line 2? **Select all that apply.**

```
1 movies |>
2   mutate(runtime = ____) |>
3   ggplot(aes(x = runtime, y = score)) +
4   geom_point(alpha = 0.5) +
5   geom_smooth(method = "lm", se = FALSE)
6   # additional code for annotating Blue City on the plot
```

- a. grepl(" mins", runtime)
- b. grep(" mins", runtime)
- c. str_remove(runtime, " mins")
- d. as.numeric(str_remove(runtime, " mins"))
- e. na.rm(runtime)

Question 8

Based on this model, order the three labeled movies in Figure 2 in decreasing order of the magnitude (absolute value) of their residuals.

- a. Winter Sleep > Rang De Basanti > Blue City
- b. Winter Sleep > Blue City > Rang De Basanti
- c. Rang De Basanti > Winter Sleep > Blue City
- d. Blue City > Winter Sleep > Rang De Basanti
- e. Blue City > Rang De Basanti > Winter Sleep


Question 9

The R-squared for the model visualized in Figure 2 is 31%. Which of the following is the best interpretation of this value?

- a. 31% of the variability in movie runtimes is explained by their scores.
- b. 31% of the variability in movie scores is explained by their runtime.
- c. The model accurately predicts scores of 31% of the movies in this sample.
- d. The model accurately predicts scores of 31% of all movies.
- e. The correlation between scores and runtimes of movies is 0.31.

Score vs. runtime or year

The visualizations below show the relationship between `score` and `runtime` as well as `score` and `year`, respectively. Additionally, the lines of best fit are overlaid on the visualizations.



The correlation coefficients of these relationships are calculated below, though some of the code and the output are missing. Answer all questions in this part based on the code and output shown below.

```
movies |>
  __blank_1__(
    r_score_runtime = cor(runtime, score),
    r_score_year = cor(year, score)
  )
```

```
# A tibble: 1 × 2
r_score_runtime r_score_year
```

```
1      <dbl>      <dbl>
0.434. __blank_2__
```

Question 10

Which of the following goes in __blank_1__?

- a. `summarize`
- b. `mutate`
- c. `group_by`
- d. `arrange`
- e. `filter`

Question 11

What can we say about the value that goes in __blank_2__?

- a. NA
- b. A value between 0 and 0.434.
- c. A value between 0.434 and 1.
- d. A value between 0 and -0.434.
- e. A value between -1 and -0.434.

Score vs. runtime and rating

In this part, we fit a model predicting `score` from `runtime` and `rating` (categorized as G, PG, PG-13, R, NC-17, and Not Rated), and name it `score_runtime_rating_fit`.

The model output for `score_runtime_rating_fit` is shown in Table 3. Answer all questions in this part based on Table 3.

Table 3: Regression output for `score_runtime_rating_fit`.

| term | estimate | std.error | statistic | p.value |
|-----------------|----------|-----------|-----------|---------|
| (Intercept) | 4.525 | 0.332 | 13.647 | 0.000 |
| runtime | 0.021 | 0.002 | 10.702 | 0.000 |
| ratingPG | -0.189 | 0.295 | -0.642 | 0.521 |
| ratingPG-13 | -0.452 | 0.292 | -1.547 | 0.123 |
| ratingR | -0.257 | 0.285 | -0.901 | 0.368 |
| ratingNC-17 | -0.355 | 0.486 | -0.730 | 0.466 |
| ratingNot Rated | -0.282 | 0.328 | -0.860 | 0.390 |

Question 12

Which of the following is **TRUE** about the intercept of `score_runtime_rating_fit`? Select all that are true.

- Keeping runtime constant, G-rated movies are predicted to score, on average, 4.525 points.
- Keeping runtime constant, movies without a rating are predicted to score, on average, 4.525 points.
- Movies without a rating that are 0 minutes in length are predicted to score, on average, 4.525 points.
- All else held constant, movies that are 0 minutes in length are predicted to score, on average, 4.525 points.
- G-rated movies that are 0 minutes in length are predicted to score, on average, 4.525 points.

Question 13

Which of the following is the **best** interpretation of the slope of `runtime` in `score_runtime_rating_fit`?

- a. All else held constant, as runtime increases by 1 minute, the score of the movie increases by 0.021 points.
- b. For G-rated movies, all else held constant, as runtime increases by 1 minute, the score of the movie increases by 0.021 points.
- c. All else held constant, for each additional minute of runtime, movie scores will be higher by 0.021 points on average.
- d. G-rated movies that are 0 minutes in length are predicted to score 0.021 points on average.
- e. For each higher level of rating, the movie scores go up by 0.021 points on average.

Question 14

Fill in the blank:

R-squared for `score_runtime_rating_fit` (the model predicting `score` from `runtime` and `rating`) _____ the R-squared the model `score_runtime_fit` (for predicting `score` from `runtime` alone).

- a. is less than
- b. is equal to
- c. is greater than
- d. cannot be compared (based on the information provided) to
- e. is both greater than and less than

Question 15

The model `score_runtime_rating_fit` (the model predicting `score` from `runtime` and `rating`) can be visualized as parallel lines for each level of `rating`. Which of the following is the equation of the line for R-rated movies?

- a. $\widehat{score} = (4.525 - 0.257) + 0.021 \times runtime$
- b. $score = (4.525 - 0.257) + 0.021 \times runtime$
- c. $\widehat{score} = 4.525 + (0.021 - 0.257) \times runtime$
- d. $score = 4.525 + (0.021 - 0.257) \times runtime$
- e. $\widehat{score} = (4.525 + 0.021) - 0.257 \times runtime$

Miscellaneous

Question 16

Which of the following is the definition of a regression model? Select all that apply.

- a. $\hat{y} = b_0 + b_1 X_1$
- b. $y = \beta_0 + \beta_1 X_1$
- c. $\hat{y} = \beta_0 + \beta_1 X_1 + \epsilon$
- d. $y = \beta_0 + \beta_1 X_1 + \epsilon$

Question 17

Define the term “parsimonious model”.

Building a spam filter

The data come from incoming emails in David Diez's (one of the authors of OpenIntro textbooks) Gmail account for the first three months of 2012. All personally identifiable information has been removed. The dataset is called `email` and it's in the `openintro` package.

The outcome variable is `spam`, which takes the value 1 if the email is spam, 0 otherwise.

Question 18

- a. What type of variable is `spam`? What percent of the emails are spam?
- b. What type of variable is `dollar` - number of times a dollar sign or the word “dollar” appeared in the email? Visualize and describe its distribution, supporting your description with the appropriate summary statistics.
- c. Fit a logistic regression model predicting `spam` from `dollar`. Then, display the tidy output of the model.
- d. Using this model and the `predict()` function, predict the probability the email is spam if it contains 5 dollar signs. Based on this probability, how does the model classify this email?

i Note

To obtain the predicted probability, you can set the `type` argument in `predict()` to `"prob"`.

Question 19

- a. Fit another logistic regression model predicting `spam` from `dollar`, `winner` (indicating whether “winner” appeared in the email), and `urgent_subj` (whether the word “urgent” is in the subject of the email). Then, display the tidy output of the model.
- b. Using this model and the `augment()` function, classify each email in the `email` dataset as spam or not spam. Store the resulting data frame with an appropriate name and display the data frame as well.
- c. Using your data frame from the previous part, determine, in a single pipeline, and using `count()`, the numbers of emails:
 - that are labelled as spam that are actually spam
 - that are not labelled as spam that are actually spam
 - that are labelled as spam that are actually not spam
 - that are not labelled as spam that are actually not spam

Store the resulting data frame with an appropriate name and display the data frame as well.

- d. In a single pipeline, and using `mutate()`, calculate the false positive and false negative rates. In addition to these numbers showing in your R output, you must write a sentence that explicitly states and identified the two rates.

Question 20

- a. Fit another logistic regression model predicting `spam` from `dollar` and another variable you think would be a good predictor. Provide a 1-sentence justification for why you chose this variable. Display the tidy output of the model.
- b. Using this model and the `augment()` function, classify each email in the `email` dataset as spam or not spam. Store the resulting data frame with an appropriate name and display the data frame as well.
- c. Using your data frame from the previous part, determine, in a single pipeline, and using `count()`, the numbers of emails:
 - that are labelled as spam that are actually spam
 - that are not labelled as spam that are actually spam
 - that are labelled as spam that are actually not spam
 - that are not labelled as spam that are actually not spam

Store the resulting data frame with an appropriate name and display the data frame as well.

- d. In a single pipeline, and using `mutate()`, calculate the false positive and false negative rates. In addition to these numbers showing in your R output, you must write a sentence that explicitly states and identified the two rates.
- e. Based on the false positive and false negatives rates of this model, comment, in 1-2 sentences, on which model (one from Question 19 or Question 20) is preferable and why.